


AI Agent Systems Field Manual

Automation, Multi-Agent Workflows, Tools, Security, and Governance



CUSTOMER EDITION

Practical guidance for designing governed, testable AI agent systems.

Free customer resource | YK SYSTEMS | hub.yksystems.ca

How to use this field manual

This manual is a practical reference for founders, creators, consultants, operators, and small business owners. You do not need to build every layer at once. Start with one workflow, add clear boundaries, and expand only after the first system is reliable.

Automation moves repeated work. AI agents help manage thinking work. Agent systems connect both into a governed business operating layer.

A sensible reading path

- New to agents: read the evolution ladder, workflow patterns, stack, security, and approval gates.
- Planning a build: focus on tool decisions, testing, observability, and the implementation blueprints.
- Preparing delivery: complete the worksheets and use the Deployment Toolkit templates.

Contents

1. The Evolution Ladder
2. Agentic Workflow Pattern: Single-shot
3. Agentic Workflow Pattern: Chaining
4. Agentic Workflow Pattern: Routing
5. Agentic Workflow Pattern: Orchestrator
6. Agentic Workflow Pattern: Evaluator
7. Agentic Workflow Pattern: Tool-using
8. Agentic Workflow Pattern: Parallel
9. Agentic Workflow Pattern: Autonomous loop
10. The Full AI Agent Stack
11. Agent Stack Layer: RAG / context layer
12. Agent Stack Layer: Function calling / tool layer
13. Agent Stack Layer: MCP / access layer
14. Agent Stack Layer: CLI / local control layer
15. Agent Stack Layer: Orchestration layer
16. Agent Stack Layer: Agent-to-agent coordination
17. Agent Stack Layer: Governance layer
18. Agent Stack Layer: Audit / observability layer
19. Anatomy of an Agent Workspace
20. Workspace Component: agents/
21. Workspace Component: skills/
22. Workspace Component: commands/
23. Workspace Component: hooks/
24. Workspace Component: rules/
25. Workspace Component: output-styles/
26. Workspace Component: settings/
27. Workspace Component: mcp/
28. Workspace Component: memory/
29. Workspace Component: logs/
30. Workspace Component: evals/
31. Workspace Component: workflows/
32. Tool-by-Tool Ecosystem Guide
33. Tool-by-Tool Ecosystem Guide Continued
34. Tool Decision: ChatGPT / GPT apps
35. Tool Decision: Claude / Claude Code
36. Tool Decision: Gemini

37. Tool Decision: Cursor / Windsurf
38. Tool Decision: Codex / code agents
39. Tool Decision: n8n
40. Tool Decision: Make
41. Tool Decision: Zapier
42. Tool Decision: LangGraph
43. Tool Decision: CrewAI
44. Tool Decision: Supabase
45. Tool Decision: Airtable
46. Tool Decision: Google Sheets
47. Tool Decision: Vector databases
48. Tool Decision: Stripe
49. Tool Decision: Slack / Telegram / Discord
50. RAG and Context Systems
51. Function Calling and Structured Tool Schemas
52. MCP and Access Layer Design
53. CLI Agents and Local Control
54. No-Code Agent Workflows
55. Code-First Agent Systems
56. Multi-Agent Coordination
57. Memory and State Management
58. Testing and Evals
59. Observability and Audit Logs
60. Agent Security Threat Model
61. Agent Security Threat Model Continued
62. Security Control: Prompt injection
63. Security Control: Tool poisoning
64. Security Control: Over-permissioning
65. Security Control: Credential exposure
66. Security Control: Data leakage
67. Security Control: Autonomous action risk
68. Security Control: Memory contamination
69. Security Control: Multi-agent failure chain
70. Security Control: Payment/billing risk
71. Security Control: Unsafe local execution
72. Agent Risk Classification Framework
73. Approval Gate Design
74. Kill Switch Procedure

75. Environment Separation
76. Audit Log Schema
77. Incident Response SOP
78. Real Implementation Blueprints
79. Blueprint: Missed Lead Recovery Agent - Overview
80. Blueprint: Missed Lead Recovery Agent - Architecture
81. Blueprint: Missed Lead Recovery Agent - Build Steps
82. Blueprint: Customer Support Draft Agent - Overview
83. Blueprint: Customer Support Draft Agent - Architecture
84. Blueprint: Customer Support Draft Agent - Build Steps
85. Blueprint: Weekly Reporting Agent - Overview
86. Blueprint: Weekly Reporting Agent - Architecture
87. Blueprint: Weekly Reporting Agent - Build Steps
88. Blueprint: Content Operations Agent Crew - Overview
89. Blueprint: Content Operations Agent Crew - Architecture
90. Blueprint: Content Operations Agent Crew - Build Steps
91. Blueprint: Multi-Agent Business OS - Overview
92. Blueprint: Multi-Agent Business OS - Architecture
93. Blueprint: Multi-Agent Business OS - Build Steps
94. Production-Ready Tool Decision Matrix
95. Business Impact Matrix
96. YK SYSTEMS Delivery Model
97. Implementation Path: Workflow Audit
98. Implementation Path: Client Command Center Setup
99. Implementation Path: AI Agent System Design
100. Implementation Path: Multi-Agent Operating System
101. Implementation Path: Governance and Approval Layer
102. Worksheet: Agent Role Canvas
103. Worksheet: Agent Tool Permission Sheet
104. Worksheet: Agent Memory Plan
105. Worksheet: Approval Gate Map
106. Worksheet: Agent Eval Scorecard
107. Worksheet: MCP Server Planning Sheet
108. Worksheet: RAG Knowledge Base Planner
109. Worksheet: Business OS Architecture Canvas
110. Worksheet: Client Handoff Checklist
111. Worksheet: Monthly Agent Audit

The Evolution Ladder

Example

1

Manual work

Owner manually replies and tracks tasks.

2

Automation

Form creates CRM row and reminder.

3

AI-assisted workflow

AI summarizes lead and drafts reply.

4

Single agent

Agent prepares follow-up package.

5

Tool-using agent

Agent queries CRM and creates task.

6

Multi-agent system

Manager coordinates support, reporting, and content agents.

7

Business OS

Agents + automations + dashboards + approvals operate together.

The ladder

Manual work -> workflow automation -> AI-assisted workflows -> single agents -> tool-using agents -> multi-agent systems -> agent-powered

operating systems.

Why the ladder matters

Each layer changes the risk profile. Automation follows defined rules. Agents interpret context, draft, plan, and may choose tools. Multi-agent

systems create coordination risk. Operating systems require governance, auditability, and ownership.

Implementation principle

Do not jump to autonomous agents before the data model, workflow, permissions, and review rules exist.

Production-ready systems are built in layers:

intake, records, alerts, AI assistance, approval, tool actions, and observability.

Agentic Workflow Pattern: Single-shot

Pattern definition

One prompt, one response. Best for low-risk drafting and one-off analysis.

Business example

Draft a quick follow-up email for owner review.

When to use it

Use single-shot when the task structure matches the risk and complexity. Do not force every business problem into a multi-agent design. Start

with the smallest pattern that produces reliable value.

Implementation steps

Define input and expected output.

Define tool access and permission level.

Decide whether output is draft, recommendation, task, record update, or external action.

Add approval gate if the action touches customers, money, public messaging, or production data.

Log the final result and failure path.

Risk check

The more autonomy the pattern gives, the stronger the approval and logging layer must be. Autonomous loops must include kill switch, retry

limits, action limits, and owner visibility.

Agentic Workflow Pattern: Chaining

Pattern definition

Break the task into sequential steps where one output feeds the next.

Business example

Analyze lead -> summarize -> draft reply -> create task.

When to use it

Use chaining when the task structure matches the risk and complexity. Do not force every business problem into a multi-agent design. Start

with the smallest pattern that produces reliable value.

Implementation steps

Define input and expected output.

Define tool access and permission level.

Decide whether output is draft, recommendation, task, record update, or external action.

Add approval gate if the action touches customers, money, public messaging, or production data.

Log the final result and failure path.

Risk check

The more autonomy the pattern gives, the stronger the approval and logging layer must be. Autonomous loops must include kill switch, retry

limits, action limits, and owner visibility.

Agentic Workflow Pattern: Routing

Pattern definition

Classify a request and send it to the correct workflow or agent.

Business example

Sales request vs support vs billing vs general info.

When to use it

Use routing when the task structure matches the risk and complexity. Do not force every business problem into a multi-agent design. Start

with the smallest pattern that produces reliable value.

Implementation steps

Define input and expected output.

Define tool access and permission level.

Decide whether output is draft, recommendation, task, record update, or external action.

Add approval gate if the action touches customers, money, public messaging, or production data.

Log the final result and failure path.

Risk check

The more autonomy the pattern gives, the stronger the approval and logging layer must be. Autonomous loops must include kill switch, retry

limits, action limits, and owner visibility.

Agentic Workflow Pattern: Orchestrator

Pattern definition

A manager agent assigns work to specialists and merges results.

Business example

Manager delegates research, drafting, QA, and reporting.

When to use it

Use orchestrator when the task structure matches the risk and complexity. Do not force every business problem into a multi-agent design.

Start with the smallest pattern that produces reliable value.

Implementation steps

Define input and expected output.

Define tool access and permission level.

Decide whether output is draft, recommendation, task, record update, or external action.

Add approval gate if the action touches customers, money, public messaging, or production data.

Log the final result and failure path.

Risk check

The more autonomy the pattern gives, the stronger the approval and logging layer must be. Autonomous loops must include kill switch, retry

limits, action limits, and owner visibility.

Agentic Workflow Pattern: Evaluator

Pattern definition

A reviewing agent critiques, scores, and requests revision.

Business example

Draft -> evaluate tone/risk -> revise -> final review.

When to use it

Use evaluator when the task structure matches the risk and complexity. Do not force every business problem into a multi-agent design. Start

with the smallest pattern that produces reliable value.

Implementation steps

Define input and expected output.

Define tool access and permission level.

Decide whether output is draft, recommendation, task, record update, or external action.

Add approval gate if the action touches customers, money, public messaging, or production data.

Log the final result and failure path.

Risk check

The more autonomy the pattern gives, the stronger the approval and logging layer must be. Autonomous loops must include kill switch, retry

limits, action limits, and owner visibility.

Agentic Workflow Pattern: Tool-using

Pattern definition

Agent chooses tools to query, update, create, or trigger actions.

Business example

Search CRM, draft email, create ticket, send approval alert.

When to use it

Use tool-using when the task structure matches the risk and complexity. Do not force every business problem into a multi-agent design. Start

with the smallest pattern that produces reliable value.

Implementation steps

Define input and expected output.

Define tool access and permission level.

Decide whether output is draft, recommendation, task, record update, or external action.

Add approval gate if the action touches customers, money, public messaging, or production data.

Log the final result and failure path.

Risk check

The more autonomy the pattern gives, the stronger the approval and logging layer must be. Autonomous loops must include kill switch, retry

limits, action limits, and owner visibility.

Agentic Workflow Pattern: Parallel

Pattern definition

Multiple agents work simultaneously for speed and specialization.

Business example

Research agent + content agent + visual prompt agent.

When to use it

Use parallel when the task structure matches the risk and complexity. Do not force every business problem into a multi-agent design. Start

with the smallest pattern that produces reliable value.

Implementation steps

Define input and expected output.

Define tool access and permission level.

Decide whether output is draft, recommendation, task, record update, or external action.

Add approval gate if the action touches customers, money, public messaging, or production data.

Log the final result and failure path.

Risk check

The more autonomy the pattern gives, the stronger the approval and logging layer must be. Autonomous loops must include kill switch, retry

limits, action limits, and owner visibility.

Agentic Workflow Pattern: Autonomous loop

Pattern definition

Goal -> plan -> act -> observe -> adjust, bounded by permissions and approvals.

Business example

Daily ops agent checks leads, drafts report, requests approvals.

When to use it

Use autonomous loop when the task structure matches the risk and complexity. Do not force every business problem into a multi-agent

design. Start with the smallest pattern that produces reliable value.

Implementation steps

Define input and expected output.

Define tool access and permission level.

Decide whether output is draft, recommendation, task, record update, or external action.

Add approval gate if the action touches customers, money, public messaging, or production data.

Log the final result and failure path.

Risk check

The more autonomy the pattern gives, the stronger the approval and logging layer must be. Autonomous loops must include kill switch, retry

limits, action limits, and owner visibility.

The Full AI Agent Stack

Use case

RAG / context layer

Knowledge retrieval from documents, databases, policies, files, and prior records.

Use for policy Q&A, support docs, manuals, and CRM context.

Function calling / tool layer

Structured tool schemas that let models call safe functions with defined inputs.

Use for API calls, calculations, database lookups, and form actions.

MCP / access layer

Standardized connection pattern between models and tools/data sources.

Use to expose apps, files, APIs, and business systems through governed connectors.

CLI / local control layer

Local filesystem, terminal, Git, shell scripts, code execution, and dev tooling.

Use carefully for code projects, migrations, audits, and repo operations.

Orchestration layer

State machines, routers, planners, supervisor agents, queues, and workflows.

Use to control order, retries, approvals, and handoffs.

Agent-to-agent coordination

Specialized agents working as a crew/team with clear roles.

Use for research + writing + QA + reporting + implementation workflows.

Governance layer

Permissions, risk levels, human approval, policy checks, and kill switches.

Use for every production system that touches

customers, money, or data.

Audit / observability layer

Logs, traces, eval results, execution history, and incident reports.

Use to prove what happened and diagnose failures.

Stack overview

A serious agent system is not one prompt. It is a stack of context retrieval, tool schemas, access protocols, orchestration, agent roles,

permissions, approvals, observability, and business ownership.

Stack rule

If a layer is missing, the system may still demo well but fail under real conditions. Production-ready delivery requires the stack to be explicitly

designed, not accidentally assembled.

Agent Stack Layer: RAG / context layer

Purpose

Knowledge retrieval from documents, databases, policies, files, and prior records.

Best use case

Use for policy Q&A, support docs, manuals, and CRM context.

Design questions

What data does this layer read?

What actions can this layer trigger?

Who owns it?

What failures must be logged?

What permissions are required?

What approval is needed before external action?

Agent Stack Layer: Function calling / tool layer

Purpose

Structured tool schemas that let models call safe functions with defined inputs.

Best use case

Use for API calls, calculations, database lookups, and form actions.

Design questions

What data does this layer read?

What actions can this layer trigger?

Who owns it?

What failures must be logged?

What permissions are required?

What approval is needed before external action?

Agent Stack Layer: MCP / access layer

Purpose

Standardized connection pattern between models and tools/data sources.

Best use case

Use to expose apps, files, APIs, and business systems through governed connectors.

Design questions

What data does this layer read?

What actions can this layer trigger?

Who owns it?

What failures must be logged?

What permissions are required?

What approval is needed before external action?

Agent Stack Layer: CLI / local control layer

Purpose

Local filesystem, terminal, Git, shell scripts, code execution, and dev tooling.

Best use case

Use carefully for code projects, migrations, audits, and repo operations.

Design questions

What data does this layer read?

What actions can this layer trigger?

Who owns it?

What failures must be logged?

What permissions are required?

What approval is needed before external action?

Agent Stack Layer: Orchestration layer

Purpose

State machines, routers, planners, supervisor agents, queues, and workflows.

Best use case

Use to control order, retries, approvals, and handoffs.

Design questions

What data does this layer read?

What actions can this layer trigger?

Who owns it?

What failures must be logged?

What permissions are required?

What approval is needed before external action?

Agent Stack Layer: Agent-to-agent coordination

Purpose

Specialized agents working as a crew/team with clear roles.

Best use case

Use for research + writing + QA + reporting + implementation workflows.

Design questions

What data does this layer read?

What actions can this layer trigger?

Who owns it?

What failures must be logged?

What permissions are required?

What approval is needed before external action?

Agent Stack Layer: Governance layer

Purpose

Permissions, risk levels, human approval, policy checks, and kill switches.

Best use case

Use for every production system that touches customers, money, or data.

Design questions

What data does this layer read?

What actions can this layer trigger?

Who owns it?

What failures must be logged?

What permissions are required?

What approval is needed before external action?

Agent Stack Layer: Audit / observability layer

Purpose

Logs, traces, eval results, execution history, and incident reports.

Best use case

Use to prove what happened and diagnose failures.

Design questions

What data does this layer read?

What actions can this layer trigger?

Who owns it?

What failures must be logged?

What permissions are required?

What approval is needed before external action?

Anatomy of an Agent Workspace

Overview

Folder

Purpose

agents/

Specialized workers with role, goal, constraints, tools, and output format.

skills/

Reusable capabilities or modules that agents can invoke.

commands/

Repeatable operator commands for common actions.

hooks/

Automatic triggers that fire on events such as save, completion, or failure.

rules/

Global operating policy, boundaries, and project constraints.

output-styles/

Standard response structures for reports, code, reviews, and handoffs.

settings/

Permissions, models, tool registry, allowed actions.

mcp/

External tool/server access configuration.

memory/

Raw, working, and canonical knowledge stores.

logs/

Audit history, tool calls, failures, decisions.

evals/

Quality tests, regression checks, scoring rubrics.

workflows/

Business processes and agent orchestration maps.

Why this matters

The future of agents is moving toward structured workspaces rather than isolated chats. Agent folders, skills, rules, hooks, permissions,

memory, logs, and evals make systems repeatable and governable.

Workspace principle

Every file should have a purpose. Rules define behavior. Agents define roles. Skills define reusable abilities. Hooks trigger automation. Logs

prove what happened. Eval prevents regression.

Workspace Component: agents/

Purpose

Specialized workers with role, goal, constraints, tools, and output format.

Implementation guidance

Use agents/ only when the project needs that level of structure. Avoid creating empty architecture. Start with rules, agents, logs, and

workflows, then add skills, hooks, MCP servers, memory, and evals when the system matures.

Security note

Do not commit secrets, local overrides, API keys, or private data. Use .gitignore, environment variables, secret managers, and least-privilege

accounts.

Handoff note

Document what this component does, who owns it, how it is tested, and what happens when it fails.

Workspace Component: skills/

Purpose

Reusable capabilities or modules that agents can invoke.

Implementation guidance

Use skills/ only when the project needs that level of structure. Avoid creating empty architecture. Start with rules, agents, logs, and workflows,

then add skills, hooks, MCP servers, memory, and evals when the system matures.

Security note

Do not commit secrets, local overrides, API keys, or private data. Use .gitignore, environment variables, secret managers, and least-privilege

accounts.

Handoff note

Document what this component does, who owns it, how it is tested, and what happens when it fails.

Workspace Component: commands/

Purpose

Repeatable operator commands for common actions.

Implementation guidance

Use `commands/` only when the project needs that level of structure. Avoid creating empty architecture. Start with rules, agents, logs, and

workflows, then add skills, hooks, MCP servers, memory, and evals when the system matures.

Security note

Do not commit secrets, local overrides, API keys, or private data. Use `.gitignore`, environment variables, secret managers, and least-privilege

accounts.

Handoff note

Document what this component does, who owns it, how it is tested, and what happens when it fails.

Workspace Component: hooks/

Purpose

Automatic triggers that fire on events such as save, completion, or failure.

Implementation guidance

Use hooks/ only when the project needs that level of structure. Avoid creating empty architecture. Start with rules, agents, logs, and

workflows, then add skills, hooks, MCP servers, memory, and evals when the system matures.

Security note

Do not commit secrets, local overrides, API keys, or private data. Use .gitignore, environment variables, secret managers, and least-privilege

accounts.

Handoff note

Document what this component does, who owns it, how it is tested, and what happens when it fails.

Workspace Component: rules/

Purpose

Global operating policy, boundaries, and project constraints.

Implementation guidance

Use rules/ only when the project needs that level of structure. Avoid creating empty architecture. Start with rules, agents, logs, and workflows,

then add skills, hooks, MCP servers, memory, and evals when the system matures.

Security note

Do not commit secrets, local overrides, API keys, or private data. Use .gitignore, environment variables, secret managers, and least-privilege

accounts.

Handoff note

Document what this component does, who owns it, how it is tested, and what happens when it fails.

Workspace Component: output-styles/

Purpose

Standard response structures for reports, code, reviews, and handoffs.

Implementation guidance

Use output-styles/ only when the project needs that level of structure. Avoid creating empty architecture. Start with rules, agents, logs, and

workflows, then add skills, hooks, MCP servers, memory, and evals when the system matures.

Security note

Do not commit secrets, local overrides, API keys, or private data. Use .gitignore, environment variables, secret managers, and least-privilege

accounts.

Handoff note

Document what this component does, who owns it, how it is tested, and what happens when it fails.

Workspace Component: settings/

Purpose

Permissions, models, tool registry, allowed actions.

Implementation guidance

Use settings/ only when the project needs that level of structure. Avoid creating empty architecture. Start with rules, agents, logs, and

workflows, then add skills, hooks, MCP servers, memory, and evals when the system matures.

Security note

Do not commit secrets, local overrides, API keys, or private data. Use .gitignore, environment variables, secret managers, and least-privilege

accounts.

Handoff note

Document what this component does, who owns it, how it is tested, and what happens when it fails.

Workspace Component: mcp/

Purpose

External tool/server access configuration.

Implementation guidance

Use mcp/ only when the project needs that level of structure. Avoid creating empty architecture. Start with rules, agents, logs, and workflows,

then add skills, hooks, MCP servers, memory, and evals when the system matures.

Security note

Do not commit secrets, local overrides, API keys, or private data. Use .gitignore, environment variables, secret managers, and least-privilege

accounts.

Handoff note

Document what this component does, who owns it, how it is tested, and what happens when it fails.

Workspace Component: memory/

Purpose

Raw, working, and canonical knowledge stores.

Implementation guidance

Use memory/ only when the project needs that level of structure. Avoid creating empty architecture. Start with rules, agents, logs, and

workflows, then add skills, hooks, MCP servers, memory, and evals when the system matures.

Security note

Do not commit secrets, local overrides, API keys, or private data. Use .gitignore, environment variables, secret managers, and least-privilege

accounts.

Handoff note

Document what this component does, who owns it, how it is tested, and what happens when it fails.

Workspace Component: logs/

Purpose

Audit history, tool calls, failures, decisions.

Implementation guidance

Use logs/ only when the project needs that level of structure. Avoid creating empty architecture. Start with rules, agents, logs, and workflows,

then add skills, hooks, MCP servers, memory, and evals when the system matures.

Security note

Do not commit secrets, local overrides, API keys, or private data. Use .gitignore, environment variables, secret managers, and least-privilege

accounts.

Handoff note

Document what this component does, who owns it, how it is tested, and what happens when it fails.

Workspace Component: evals/

Purpose

Quality tests, regression checks, scoring rubrics.

Implementation guidance

Use evals/ only when the project needs that level of structure. Avoid creating empty architecture. Start with rules, agents, logs, and workflows,

then add skills, hooks, MCP servers, memory, and evals when the system matures.

Security note

Do not commit secrets, local overrides, API keys, or private data. Use .gitignore, environment variables, secret managers, and least-privilege

accounts.

Handoff note

Document what this component does, who owns it, how it is tested, and what happens when it fails.

Workspace Component: workflows/

Purpose

Business processes and agent orchestration maps.

Implementation guidance

Use workflows/ only when the project needs that level of structure. Avoid creating empty architecture. Start with rules, agents, logs, and

workflows, then add skills, hooks, MCP servers, memory, and evals when the system matures.

Security note

Do not commit secrets, local overrides, API keys, or private data. Use .gitignore, environment variables, secret managers, and least-privilege

accounts.

Handoff note

Document what this component does, who owns it, how it is tested, and what happens when it fails.

Tool-by-Tool Ecosystem Guide

Overview

Tool

Caution

ChatGPT / GPT apps

Assistant workflows, research, drafting,
analysis, no-code ideation

Broad reasoning, tool ecosystems,
GPTs/assistants, fast user-facing

Do not let chat-only workflows become
unlogged production actions.

Claude / Claude Code

Deep writing, codebase navigation,
coding agents, context-heavy work

Large context workflows, coding, review,
project understanding

Require repo permissions and shell
actions to be scoped carefully.

Gemini

Google ecosystem, long context,
multimodal workflows

Workspace-adjacent teams, docs, sheets,
search context

Check tool access, privacy, and
enterprise settings.

Cursor / Windsurf

IDE agentic coding workflows

Multi-file coding, refactors, debugging,
developer velocity

Use branch isolation, tests, and PR
review.

Codex / code agents

Autonomous code task execution and
PR workflows

Repo tasks, tests, code changes,

issue-driven work

Require scoped tasks and approval for destructive actions.

n8n

Controlled automation engine, webhooks, APIs, branching

Advanced business workflows, self-hosting, custom logic

Requires workflow discipline and credential hygiene.

Make

Visual no-code automation with routers and integrations

Business users who need visible scenario logic

Can get messy without naming and documentation.

Zapier

Fast app-to-app automation

Simple triggers, form to CRM, email, alerts

Less suited for complex stateful agent orchestration.

LangGraph

Stateful agent graphs and controllable workflows

Multi-step, stateful, durable, multi-agent systems

Requires engineering ability and careful state design.

CrewAI

Role-based agents, tasks, crews, flows
Specialized agent teams and business process prototypes

Needs strong task definitions and monitoring.

How to choose tools

Tool choice should follow job-to-be-done, risk level, operator skill, data sensitivity, integrations, budget, and maintenance capacity. Do not

choose tools because they sound advanced.

Decision model

Beginner systems prioritize ease and clarity. Intermediate systems prioritize structured workflows and CRM visibility.

Advanced systems

prioritize state, governance, observability, security, and custom integration.

Tool-by-Tool Ecosystem Guide Continued

Overview

Tool

Caution

Supabase

Postgres backend, auth, storage, edge functions

Production CRM, logs, dashboards, SaaS foundations

RLS and secret management are non-negotiable.

Airtable

Structured no-code database and views

Small business CRMs and operational views

Not ideal as high-volume backend.

Google Sheets

Fast, simple CRM/log/testing data

Beginner systems and exports

Weak permissions and relational control.

Vector databases

Semantic retrieval and memory

RAG over docs, support, policies, manuals

Bad chunking or stale docs produce weak answers.

Stripe

Payments and billing events

Payment-triggered onboarding/delivery

Verify event IDs, status, amount, currency, product IDs.

Slack / Telegram /

Discord

Alerts, approvals, team notifications

Owner alerts and quick review queues

Avoid sending sensitive data to

informal channels.

Operational note

The right stack is rarely one tool. A real system may combine a form, CRM, automation engine, AI model, database, alert channel, approval

queue, logging layer, and service handoff.

Upgrade path

Start with visible workflow value. Then add agents where judgment, drafting, research, or tool coordination creates measurable operational

improvement.

Tool Decision: ChatGPT / GPT apps

Best role

Assistant workflows, research, drafting, analysis, no-code ideation

Strength

Broad reasoning, tool ecosystems, GPTs/assistants, fast user-facing workflows

Caution

Do not let chat-only workflows become unlogged production actions.

Use when

Use ChatGPT / GPT apps when its strengths match the actual workflow requirement. If the business need can be met with a simpler tool, do

not add unnecessary complexity.

Do not use when

Avoid this direction when the operator cannot maintain it, the data model is unclear, or the governance requirements are not defined.

Tool Decision: Claude / Claude Code

Best role

Deep writing, codebase navigation, coding agents, context-heavy work

Strength

Large context workflows, coding, review, project understanding

Caution

Require repo permissions and shell actions to be scoped carefully.

Use when

Use Claude / Claude Code when its strengths match the actual workflow requirement. If the business need can be met with a simpler tool, do

not add unnecessary complexity.

Do not use when

Avoid this direction when the operator cannot maintain it, the data model is unclear, or the governance requirements are not defined.

Tool Decision: Gemini

Best role

Google ecosystem, long context, multimodal workflows

Strength

Workspace-adjacent teams, docs, sheets, search context

Caution

Check tool access, privacy, and enterprise settings.

Use when

Use Gemini when its strengths match the actual workflow requirement. If the business need can be met with a simpler tool, do not add

unnecessary complexity.

Do not use when

Avoid this direction when the operator cannot maintain it, the data model is unclear, or the governance requirements are not defined.

Tool Decision: Cursor / Windsurf

Best role

IDE agentic coding workflows

Strength

Multi-file coding, refactors, debugging, developer velocity

Caution

Use branch isolation, tests, and PR review.

Use when

Use Cursor / Windsurf when its strengths match the actual workflow requirement. If the business need can be met with a simpler tool, do not

add unnecessary complexity.

Do not use when

Avoid this direction when the operator cannot maintain it, the data model is unclear, or the governance requirements are not defined.

Tool Decision: Codex / code agents

Best role

Autonomous code task execution and PR workflows

Strength

Repo tasks, tests, code changes, issue-driven work

Caution

Require scoped tasks and approval for destructive actions.

Use when

Use Codex / code agents when its strengths match the actual workflow requirement. If the business need can be met with a simpler tool, do

not add unnecessary complexity.

Do not use when

Avoid this direction when the operator cannot maintain it, the data model is unclear, or the governance requirements are not defined.

Tool Decision: n8n

Best role

Controlled automation engine, webhooks, APIs, branching

Strength

Advanced business workflows, self-hosting, custom logic

Caution

Requires workflow discipline and credential hygiene.

Use when

Use n8n when its strengths match the actual workflow requirement. If the business need can be met with a simpler tool, do not add

unnecessary complexity.

Do not use when

Avoid this direction when the operator cannot maintain it, the data model is unclear, or the governance requirements are not defined.

Tool Decision: Make

Best role

Visual no-code automation with routers and integrations

Strength

Business users who need visible scenario logic

Caution

Can get messy without naming and documentation.

Use when

Use Make when its strengths match the actual workflow requirement. If the business need can be met with a simpler tool, do not add

unnecessary complexity.

Do not use when

Avoid this direction when the operator cannot maintain it, the data model is unclear, or the governance requirements are not defined.

Tool Decision: Zapier

Best role

Fast app-to-app automation

Strength

Simple triggers, form to CRM, email, alerts

Caution

Less suited for complex stateful agent orchestration.

Use when

Use Zapier when its strengths match the actual workflow requirement. If the business need can be met with a simpler tool, do not add

unnecessary complexity.

Do not use when

Avoid this direction when the operator cannot maintain it, the data model is unclear, or the governance requirements are not defined.

Tool Decision: LangGraph

Best role

Stateful agent graphs and controllable workflows

Strength

Multi-step, stateful, durable, multi-agent systems

Caution

Requires engineering ability and careful state design.

Use when

Use LangGraph when its strengths match the actual workflow requirement. If the business need can be met with a simpler tool, do not add

unnecessary complexity.

Do not use when

Avoid this direction when the operator cannot maintain it, the data model is unclear, or the governance requirements are not defined.

Tool Decision: CrewAI

Best role

Role-based agents, tasks, crews, flows

Strength

Specialized agent teams and business process prototypes

Caution

Needs strong task definitions and monitoring.

Use when

Use CrewAI when its strengths match the actual workflow requirement. If the business need can be met with a simpler tool, do not add

unnecessary complexity.

Do not use when

Avoid this direction when the operator cannot maintain it, the data model is unclear, or the governance requirements are not defined.

Tool Decision: Supabase

Best role

Postgres backend, auth, storage, edge functions

Strength

Production CRM, logs, dashboards, SaaS foundations

Caution

RLS and secret management are non-negotiable.

Use when

Use Supabase when its strengths match the actual workflow requirement. If the business need can be met with a simpler tool, do not add

unnecessary complexity.

Do not use when

Avoid this direction when the operator cannot maintain it, the data model is unclear, or the governance requirements are not defined.

Tool Decision: Airtable

Best role

Structured no-code database and views

Strength

Small business CRMs and operational views

Caution

Not ideal as high-volume backend.

Use when

Use Airtable when its strengths match the actual workflow requirement. If the business need can be met with a simpler tool, do not add

unnecessary complexity.

Do not use when

Avoid this direction when the operator cannot maintain it, the data model is unclear, or the governance requirements are not defined.

Tool Decision: Google Sheets

Best role

Fast, simple CRM/log/testing data

Strength

Beginner systems and exports

Caution

Weak permissions and relational control.

Use when

Use Google Sheets when its strengths match the actual workflow requirement. If the business need can be met with a simpler tool, do not add

unnecessary complexity.

Do not use when

Avoid this direction when the operator cannot maintain it, the data model is unclear, or the governance requirements are not defined.

Tool Decision: Vector databases

Best role

Semantic retrieval and memory

Strength

RAG over docs, support, policies, manuals

Caution

Bad chunking or stale docs produce weak answers.

Use when

Use Vector databases when its strengths match the actual workflow requirement. If the business need can be met with a simpler tool, do not

add unnecessary complexity.

Do not use when

Avoid this direction when the operator cannot maintain it, the data model is unclear, or the governance requirements are not defined.

Tool Decision: Stripe

Best role

Payments and billing events

Strength

Payment-triggered onboarding/delivery

Caution

Verify event IDs, status, amount, currency, product IDs.

Use when

Use Stripe when its strengths match the actual workflow requirement. If the business need can be met with a simpler tool, do not add

unnecessary complexity.

Do not use when

Avoid this direction when the operator cannot maintain it, the data model is unclear, or the governance requirements are not defined.

Tool Decision: Slack / Telegram / Discord

Best role

Alerts, approvals, team notifications

Strength

Owner alerts and quick review queues

Caution

Avoid sending sensitive data to informal channels.

Use when

Use Slack / Telegram / Discord when its strengths match the actual workflow requirement. If the business need can be met with a simpler tool,

do not add unnecessary complexity.

Do not use when

Avoid this direction when the operator cannot maintain it, the data model is unclear, or the governance requirements are not defined.

RAG and Context Systems

Explanation

RAG connects agents to external knowledge such as PDFs, docs, support articles, policies, CRM records, and manuals. Good RAG depends on

source quality, chunking strategy, metadata, retrieval filters, freshness, and citation/traceability.

Implementation design questions

What input does this module require?

What output should it produce?

Which tool or database does it touch?

What can go wrong?

What must be logged?

What requires approval?

Reliable standard

This module should be tested with happy path, missing data, malformed data, denied permission, and failure cases before production.

Function Calling and Structured Tool Schemas

Explanation

Function calling turns model intent into structured operations. Strong tool schemas define exact inputs, types, constraints, error responses,

dry-run behavior, and allowed scopes.

Implementation design questions

What input does this module require?

What output should it produce?

Which tool or database does it touch?

What can go wrong?

What must be logged?

What requires approval?

Reliable standard

This module should be tested with happy path, missing data, malformed data, denied permission, and failure cases before production.

MCP and Access Layer Design

Explanation

MCP-style systems standardize agent access to tools and data sources. Treat MCP servers like privileged infrastructure: review manifests,

permissions, transport, authentication, and tool descriptions.

Implementation design questions

What input does this module require?

What output should it produce?

Which tool or database does it touch?

What can go wrong?

What must be logged?

What requires approval?

Reliable standard

This module should be tested with happy path, missing data, malformed data, denied permission, and failure cases before production.

CLI Agents and Local Control

Explanation

CLI agents can modify files, run tests, use Git, call shell scripts, and operate codebases. This is powerful and risky. Use branches, sandboxes,

tests, command allowlists, and approval gates.

Implementation design questions

What input does this module require?

What output should it produce?

Which tool or database does it touch?

What can go wrong?

What must be logged?

What requires approval?

Reliable standard

This module should be tested with happy path, missing data, malformed data, denied permission, and failure cases before production.

No-Code Agent Workflows

Explanation

No-code agent systems combine forms, CRMs, automation engines, AI modules, alerts, emails, and approvals. They are easier to maintain for

non-engineering teams but still require governance.

Implementation design questions

What input does this module require?

What output should it produce?

Which tool or database does it touch?

What can go wrong?

What must be logged?

What requires approval?

Reliable standard

This module should be tested with happy path, missing data, malformed data, denied permission, and failure cases before production.

Code-First Agent Systems

Explanation

Code-first systems enable state machines, typed schemas, custom APIs, evals, durable execution, observability, and complex multi-agent

coordination. They require engineering discipline and deployment practices.

Implementation design questions

What input does this module require?

What output should it produce?

Which tool or database does it touch?

What can go wrong?

What must be logged?

What requires approval?

Reliable standard

This module should be tested with happy path, missing data, malformed data, denied permission, and failure cases before production.

Multi-Agent Coordination

Explanation

Multi-agent systems split work between manager, researcher, operator, reviewer, support, reporting, and implementation agents. Coordination

must define roles, handoffs, conflict resolution, and final authority.

Implementation design questions

What input does this module require?

What output should it produce?

Which tool or database does it touch?

What can go wrong?

What must be logged?

What requires approval?

Reliable standard

This module should be tested with happy path, missing data, malformed data, denied permission, and failure cases before production.

Memory and State Management

Explanation

Agents need context, but memory creates risk. Separate raw memory, working memory, and canonical memory. Add review, source labels,

expiry rules, and deletion procedures.

Implementation design questions

What input does this module require?

What output should it produce?

Which tool or database does it touch?

What can go wrong?

What must be logged?

What requires approval?

Reliable standard

This module should be tested with happy path, missing data, malformed data, denied permission, and failure cases before production.

Testing and Evals

Explanation

Agent systems need tests for correctness, safety, tool use, latency, cost, consistency, regression, and risk. Evals should include normal cases,

edge cases, adversarial prompts, and tool failures.

Implementation design questions

What input does this module require?

What output should it produce?

Which tool or database does it touch?

What can go wrong?

What must be logged?

What requires approval?

Reliable standard

This module should be tested with happy path, missing data, malformed data, denied permission, and failure cases before production.

Observability and Audit Logs

Explanation

Logs and traces show what the agent saw, decided, called, returned, and changed. Production-ready systems need observable runs, approval history,

error logs, and incident review procedures.

Implementation design questions

What input does this module require?

What output should it produce?

Which tool or database does it touch?

What can go wrong?

What must be logged?

What requires approval?

Reliable standard

This module should be tested with happy path, missing data, malformed data, denied permission, and failure cases before production.

Agent Security Threat Model

Threat

Description

Control

Prompt injection

Malicious or hidden instructions alter agent behavior.

Separate user content from instructions, add tool-call policy checks, detect suspicious instructions.

Tool poisoning

Tool descriptions or schemas contain hidden adversarial instructions.

Review tool manifests, pin versions, sign manifests, run semantic checks.

Over-permissioning

Agent has more tool access than needed.

Least privilege, read-only first, separate high-risk actions behind approval.

Credential exposure

Secrets leak through logs, prompts, repos, or outputs.

Secret managers, environment variables, redaction, no secrets in prompts/docs.

Data leakage

Private records are sent to unauthorized tools or channels.

Data classification, access scopes, DLP checks, channel limits.

Why security changes with agents

Agents are different from chatbots because they may select tools, call APIs, update records, read private files, execute code, or trigger

workflows. This creates new risk surfaces.

Security operating principle

Read-only before write. Draft before send. Dry-run before execute. Approval before external action. Logs before trust. Kill switch before

autonomy.

Agent Security Threat Model Continued

Threat

Description

Control

Autonomous action risk

Agent acts externally without confirmation.

Approval gates, dry-run mode, action queues, kill switch.

Memory contamination

Bad or outdated data pollutes future decisions.

Memory review, TTLs, source labels, separate raw/work/canon memory.

Multi-agent failure chain

One weak agent output corrupts another agent action.

Typed handoffs, evaluator agents, confidence thresholds, human review.

Payment/billing risk

Incorrect billing, refund, subscription, or access changes.

Verify events, idempotency, audit logs, manual approval for refunds.

Unsafe local execution

CLI/shell/file actions damage systems.

Sandboxing, branch isolation, command blocklist, dry-run, approvals.

Security note

Prompt injection and tool poisoning deserve special attention because they can hide inside user messages, documents, tool descriptions, web

pages, or retrieved context. The defense is layered, not one prompt line.

Control philosophy

Treat agent tools as permissions, not conveniences. Each permission increases potential blast radius.

Security Control: Prompt injection

Threat

Malicious or hidden instructions alter agent behavior.

Primary control

Separate user content from instructions, add tool-call policy checks, detect suspicious instructions.

Implementation checklist

Classify the risk level.

Reduce permissions.

Add dry-run if possible.

Add human approval for external or irreversible actions.

Log input, tool call, decision, and result.

Add incident response path.

Service delivery note

Production-ready clients will trust agent systems more when security is operational, not just mentioned. Deliver permission maps, approval rules, and incident SOPs.

Security Control: Tool poisoning

Threat

Tool descriptions or schemas contain hidden adversarial instructions.

Primary control

Review tool manifests, pin versions, sign manifests, run semantic checks.

Implementation checklist

Classify the risk level.

Reduce permissions.

Add dry-run if possible.

Add human approval for external or irreversible actions.

Log input, tool call, decision, and result.

Add incident response path.

Service delivery note

Production-ready clients will trust agent systems more when security is operational, not just mentioned. Deliver permission maps, approval rules, and incident SOPs.

Security Control: Over-permissioning

Threat

Agent has more tool access than needed.

Primary control

Least privilege, read-only first, separate high-risk actions behind approval.

Implementation checklist

Classify the risk level.

Reduce permissions.

Add dry-run if possible.

Add human approval for external or irreversible actions.

Log input, tool call, decision, and result.

Add incident response path.

Service delivery note

Production-ready clients will trust agent systems more when security is operational, not just mentioned. Deliver permission maps, approval rules, and incident SOPs.

Security Control: Credential exposure

Threat

Secrets leak through logs, prompts, repos, or outputs.

Primary control

Secret managers, environment variables, redaction, no secrets in prompts/docs.

Implementation checklist

Classify the risk level.

Reduce permissions.

Add dry-run if possible.

Add human approval for external or irreversible actions.

Log input, tool call, decision, and result.

Add incident response path.

Service delivery note

Production-ready clients will trust agent systems more when security is operational, not just mentioned. Deliver permission maps, approval rules, and incident SOPs.

Security Control: Data leakage

Threat

Private records are sent to unauthorized tools or channels.

Primary control

Data classification, access scopes, DLP checks, channel limits.

Implementation checklist

Classify the risk level.

Reduce permissions.

Add dry-run if possible.

Add human approval for external or irreversible actions.

Log input, tool call, decision, and result.

Add incident response path.

Service delivery note

Production-ready clients will trust agent systems more when security is operational, not just mentioned. Deliver permission maps, approval rules, and incident SOPs.

Security Control: Autonomous action risk

Threat

Agent acts externally without confirmation.

Primary control

Approval gates, dry-run mode, action queues, kill switch.

Implementation checklist

Classify the risk level.

Reduce permissions.

Add dry-run if possible.

Add human approval for external or irreversible actions.

Log input, tool call, decision, and result.

Add incident response path.

Service delivery note

Production-ready clients will trust agent systems more when security is operational, not just mentioned. Deliver permission maps, approval rules, and incident SOPs.

Security Control: Memory contamination

Threat

Bad or outdated data pollutes future decisions.

Primary control

Memory review, TTLs, source labels, separate raw/work/canon memory.

Implementation checklist

Classify the risk level.

Reduce permissions.

Add dry-run if possible.

Add human approval for external or irreversible actions.

Log input, tool call, decision, and result.

Add incident response path.

Service delivery note

Production-ready clients will trust agent systems more when security is operational, not just mentioned. Deliver permission maps, approval rules, and incident SOPs.

Security Control: Multi-agent failure chain

Threat

One weak agent output corrupts another agent action.

Primary control

Typed handoffs, evaluator agents, confidence thresholds, human review.

Implementation checklist

Classify the risk level.

Reduce permissions.

Add dry-run if possible.

Add human approval for external or irreversible actions.

Log input, tool call, decision, and result.

Add incident response path.

Service delivery note

Production-ready clients will trust agent systems more when security is operational, not just mentioned. Deliver permission maps, approval rules, and incident SOPs.

Security Control: Payment/billing risk

Threat

Incorrect billing, refund, subscription, or access changes.

Primary control

Verify events, idempotency, audit logs, manual approval for refunds.

Implementation checklist

Classify the risk level.

Reduce permissions.

Add dry-run if possible.

Add human approval for external or irreversible actions.

Log input, tool call, decision, and result.

Add incident response path.

Service delivery note

Production-ready clients will trust agent systems more when security is operational, not just mentioned. Deliver permission maps, approval rules, and incident SOPs.

Security Control: Unsafe local execution

Threat

CLI/shell/file actions damage systems.

Primary control

Sandboxing, branch isolation, command blocklist, dry-run, approvals.

Implementation checklist

Classify the risk level.

Reduce permissions.

Add dry-run if possible.

Add human approval for external or irreversible actions.

Log input, tool call, decision, and result.

Add incident response path.

Service delivery note

Production-ready clients will trust agent systems more when security is operational, not just mentioned. Deliver permission maps, approval rules, and incident SOPs.

Agent Risk Classification Framework

Overview

Level

Permission

Examples

0

Read-only summarize

Summarize docs, tickets, CRM records.

1

Draft only

Draft emails, reports, messages, tasks.

2

Prepare internal records

Create task draft, prep CRM update, create approval request.

3

Low-risk internal update

Update tags, add internal note, mark reviewed.

4

External action with approval

Send email, update customer, trigger onboarding.

5

Human-only

Refund, delete, legal, production, public statement, billing changes.

Purpose

Risk classification tells the builder what the agent may do automatically, what requires approval, and what must remain human-controlled.

Rule

Risk level should be assigned per action, not per agent. A support agent may safely summarize tickets but require approval for refunds or

policy exceptions.

Approval Gate Design

Overview

Operating explanation

Approval gates should capture proposed action, reason, risk, affected record, draft/output, approver, decision, timestamp, and final execution

result.

Required fields

Owner

Trigger

Action type

Risk level

Approval status

Tool used

Record affected

Result

Error/failure details

Next review date

Reliable standard

For production-ready delivery, governance must be visible in the handoff packet. It should not live only inside the builder's memory.

Kill Switch Procedure

Overview

Operating explanation

A kill switch disables external actions, tool calls, scheduled runs, or agent access quickly when behavior is unsafe, unclear, or under incident

review.

Required fields

Owner

Trigger

Action type

Risk level

Approval status

Tool used

Record affected

Result

Error/failure details

Next review date

Reliable standard

For production-ready delivery, governance must be visible in the handoff packet. It should not live only inside the builder’s memory.

Environment Separation

Overview

Operating explanation

Use sandbox/test data before production. Keep dev, staging, and production credentials separate. Never test destructive actions on live

records.

Required fields

Owner

Trigger

Action type

Risk level

Approval status

Tool used

Record affected

Result

Error/failure details

Next review date

Reliable standard

For production-ready delivery, governance must be visible in the handoff packet. It should not live only inside the builder’s memory.

Audit Log Schema

Overview

Operating explanation

Audit logs should include timestamp, run ID, agent, tool, input summary, action, result, approval status, affected record, error, and owner.

Required fields

Owner

Trigger

Action type

Risk level

Approval status

Tool used

Record affected

Result

Error/failure details

Next review date

Reliable standard

For production-ready delivery, governance must be visible in the handoff packet. It should not live only inside the builder’s memory.

Incident Response SOP

Overview

Operating explanation

Incident response should define detection, pause, scope, preserve logs, notify owner, repair, test, restore, and post-incident improvement.

Required fields

Owner

Trigger

Action type

Risk level

Approval status

Tool used

Record affected

Result

Error/failure details

Next review date

Reliable standard

For production-ready delivery, governance must be visible in the handoff packet. It should not live only inside the builder's memory.

Real Implementation Blueprints

Why this section exists

Production-ready teams need proof of implementation thinking. They need to see exact systems, not only concepts. Each blueprint includes purpose,

inputs, agent roles, tools, data flow, approval gates, logs, failure paths, build steps, tests, and deliverables.

Blueprint standard

Every implementation blueprint should be understandable without the automation editor. The architecture must make the business value and

risk controls visible.

Blueprint: Missed Lead Recovery Agent - Overview

Purpose

Detects leads that never received follow-up and prepares owner-approved recovery messages.

Inputs

CRM leads table
status and last_contacted_at
lead source
owner notes
email history

Agent actions

Read CRM
find stale leads
summarize context
draft recovery email
send approval alert
log decision

Actions requiring approval

sending email
offering discount
changing lead status
marking lead lost

Client deliverables

lead recovery queue
draft email pack
owner alert
CRM log
weekly missed-lead report

Blueprint: Missed Lead Recovery Agent - Architecture

Data flow

Input sources feed a central record. The agent reads the record, retrieves context if needed, prepares a draft/recommendation/action package,

sends it to an approval queue when required, then logs outcome after owner decision.

Tool stack recommendation

Use a form or trigger for intake, Airtable/Sheets/Supabase for records, n8n/Make/Zapier for workflow routing, AI model for analysis/drafting,

Gmail/Slack/Telegram for communication, and an audit log for every major event.

Failure paths

Missing required input -> invalid/review queue.

AI output invalid -> fallback summary and owner alert.

Tool failure -> error log and no false success.

Approval expired -> reminder and no external action.

Duplicate record -> flag or merge based on rule.

Blueprint: Missed Lead Recovery Agent - Build Steps

Overview

Build sequence

Define business objective and success metric.

Create data model and required statuses.

Build intake/trigger.

Add validation and normalization.

Add agent analysis/draft step.

Add approval gate for risky outputs.

Add CRM/task/log updates.

Add owner alert.

Add edge-case tests.

Create handoff SOP.

Testing checklist

Complete happy path.

Missing required field.

Low-quality input.

Duplicate input.

Tool unavailable.

Approval denied.

Approval accepted.

Logging visible.

Manual fallback available.

Business impact

This blueprint should reduce missed work, make priorities visible, create faster review loops, and preserve human control where risk matters.

Blueprint: Customer Support Draft Agent - Overview

Purpose

Classifies support requests and drafts answers using approved knowledge before human review.

Inputs

support form

customer account

order ID

knowledge base

policy docs

Agent actions

classify issue

retrieve answer

draft response

flag policy/risk

create ticket note

Actions requiring approval

refunds

policy exceptions

legal claims

public incident statements

Client deliverables

ticket classification

draft response

risk flag

knowledge sources used

handoff note

Blueprint: Customer Support Draft Agent - Architecture

Data flow

Input sources feed a central record. The agent reads the record, retrieves context if needed, prepares a draft/recommendation/action package,

sends it to an approval queue when required, then logs outcome after owner decision.

Tool stack recommendation

Use a form or trigger for intake, Airtable/Sheets/Supabase for records, n8n/Make/Zapier for workflow routing, AI model for analysis/drafting,

Gmail/Slack/Telegram for communication, and an audit log for every major event.

Failure paths

Missing required input -> invalid/review queue.

AI output invalid -> fallback summary and owner alert.

Tool failure -> error log and no false success.

Approval expired -> reminder and no external action.

Duplicate record -> flag or merge based on rule.

Blueprint: Customer Support Draft Agent - Build Steps

Overview

Build sequence

Define business objective and success metric.

Create data model and required statuses.

Build intake/trigger.

Add validation and normalization.

Add agent analysis/draft step.

Add approval gate for risky outputs.

Add CRM/task/log updates.

Add owner alert.

Add edge-case tests.

Create handoff SOP.

Testing checklist

Complete happy path.

Missing required field.

Low-quality input.

Duplicate input.

Tool unavailable.

Approval denied.

Approval accepted.

Logging visible.

Manual fallback available.

Business impact

This blueprint should reduce missed work, make priorities visible, create faster review loops, and preserve human control where risk matters.

Blueprint: Weekly Reporting Agent - Overview

Purpose

Summarizes weekly business operations from CRM/sheets and highlights missed actions.

Inputs

CRM records

payments

support tickets

tasks

content log

Agent actions

aggregate metrics

find stalled leads

summarize wins/problems

recommend next actions

Actions requiring approval

external sending

financial changes

staff performance judgments

Client deliverables

weekly report

missed follow-up list

risk list

next-action plan

Blueprint: Weekly Reporting Agent - Architecture

Data flow

Input sources feed a central record. The agent reads the record, retrieves context if needed, prepares a draft/recommendation/action package,

sends it to an approval queue when required, then logs outcome after owner decision.

Tool stack recommendation

Use a form or trigger for intake, Airtable/Sheets/Supabase for records, n8n/Make/Zapier for workflow routing, AI model for analysis/drafting,

Gmail/Slack/Telegram for communication, and an audit log for every major event.

Failure paths

Missing required input -> invalid/review queue.

AI output invalid -> fallback summary and owner alert.

Tool failure -> error log and no false success.

Approval expired -> reminder and no external action.

Duplicate record -> flag or merge based on rule.

Blueprint: Weekly Reporting Agent - Build Steps

Overview

Build sequence

Define business objective and success metric.

Create data model and required statuses.

Build intake/trigger.

Add validation and normalization.

Add agent analysis/draft step.

Add approval gate for risky outputs.

Add CRM/task/log updates.

Add owner alert.

Add edge-case tests.

Create handoff SOP.

Testing checklist

Complete happy path.

Missing required field.

Low-quality input.

Duplicate input.

Tool unavailable.

Approval denied.

Approval accepted.

Logging visible.

Manual fallback available.

Business impact

This blueprint should reduce missed work, make priorities visible, create faster review loops, and preserve human control where risk matters.

Blueprint: Content Operations Agent Crew - Overview

Purpose

Coordinates research, writing, visual prompts, review, and scheduling suggestions.

Inputs

content ideas

brand rules

audience profile

offer list

prior posts

Agent actions

research

outline

write

generate visual prompts

review for tone/risk

prepare schedule

Actions requiring approval

publishing

claims

pricing offers

controversial content

Client deliverables

content brief

caption pack

visual prompt pack

review notes

scheduler-ready draft

Blueprint: Content Operations Agent Crew - Architecture

Data flow

Input sources feed a central record. The agent reads the record, retrieves context if needed, prepares a draft/recommendation/action package,

sends it to an approval queue when required, then logs outcome after owner decision.

Tool stack recommendation

Use a form or trigger for intake, Airtable/Sheets/Supabase for records, n8n/Make/Zapier for workflow routing, AI model for analysis/drafting,

Gmail/Slack/Telegram for communication, and an audit log for every major event.

Failure paths

Missing required input -> invalid/review queue.

AI output invalid -> fallback summary and owner alert.

Tool failure -> error log and no false success.

Approval expired -> reminder and no external action.

Duplicate record -> flag or merge based on rule.

Blueprint: Content Operations Agent Crew - Build Steps

Overview

Build sequence

Define business objective and success metric.

Create data model and required statuses.

Build intake/trigger.

Add validation and normalization.

Add agent analysis/draft step.

Add approval gate for risky outputs.

Add CRM/task/log updates.

Add owner alert.

Add edge-case tests.

Create handoff SOP.

Testing checklist

Complete happy path.

Missing required field.

Low-quality input.

Duplicate input.

Tool unavailable.

Approval denied.

Approval accepted.

Logging visible.

Manual fallback available.

Business impact

This blueprint should reduce missed work, make priorities visible, create faster review loops, and preserve human control where risk matters.

Blueprint: Multi-Agent Business OS - Overview

Purpose

Manager agent coordinates specialist agents across operations, reporting, CRM, support, and content.

Inputs

business goals

CRM

inbox/tickets

docs

calendar/tasks

approval rules

Agent actions

plan day

delegate tasks

collect outputs

evaluate results

request approvals

log actions

Actions requiring approval

external customer actions

payments

deletions

legal language

production system changes

Client deliverables

daily command report

approval queue

agent outputs

audit log

owner dashboard

Blueprint: Multi-Agent Business OS - Architecture

Data flow

Input sources feed a central record. The agent reads the record, retrieves context if needed, prepares a draft/recommendation/action package,

sends it to an approval queue when required, then logs outcome after owner decision.

Tool stack recommendation

Use a form or trigger for intake, Airtable/Sheets/Supabase for records, n8n/Make/Zapier for workflow routing, AI model for analysis/drafting,

Gmail/Slack/Telegram for communication, and an audit log for every major event.

Failure paths

Missing required input -> invalid/review queue.

AI output invalid -> fallback summary and owner alert.

Tool failure -> error log and no false success.

Approval expired -> reminder and no external action.

Duplicate record -> flag or merge based on rule.

Blueprint: Multi-Agent Business OS - Build Steps

Overview

Build sequence

Define business objective and success metric.

Create data model and required statuses.

Build intake/trigger.

Add validation and normalization.

Add agent analysis/draft step.

Add approval gate for risky outputs.

Add CRM/task/log updates.

Add owner alert.

Add edge-case tests.

Create handoff SOP.

Testing checklist

Complete happy path.

Missing required field.

Low-quality input.

Duplicate input.

Tool unavailable.

Approval denied.

Approval accepted.

Logging visible.

Manual fallback available.

Business impact

This blueprint should reduce missed work, make priorities visible, create faster review loops, and preserve human control where risk matters.

Production-Ready Tool Decision Matrix

Overview

Need

Best direction

Why

Beginner resource delivery

Gumroad / Drive / simple hub

Fast distribution and low friction.

Basic lead capture

Form + Sheets/Airtable + email

Clear, easy to maintain.

Controlled workflow automation

n8n

Branching, APIs, webhooks, self-hosting potential.

Visual no-code routing

Make

Good for operators who need visible scenarios.

Simple app connections

Zapier

Fastest beginner path.

Production database

Supabase

Postgres, auth, storage, logs, dashboards.

Stateful agent workflows

LangGraph

Graphs, state, handoffs, control.

Role-based multi-agent

CrewAI

Agents, tasks, crews, flows.

Coding agent workflow

Claude Code / Cursor / Codex

Repo work, code review, tests, PRs.

Agent tool access

MCP / APIs / webhooks

Standardized tool and data access.

Knowledge Q&A;

RAG + vector DB

Retrieval from docs and knowledge bases.

High-risk operations

approval queue + audit logs

Human control and traceability.

How to use this matrix

Use the matrix to choose the simplest stack that satisfies the business need and risk profile. The reliable move is not choosing the fanciest

tool. The reliable move is choosing the tool that can be maintained, governed, and trusted.

Business Impact Matrix

Overview

System

Saves time by

Reduces risk by

Improves revenue by

Lead Follow-Up Agent

Drafting and prioritizing replies

Flagging missed leads

Faster response to inquiries

CRM Cleanup Agent

Organizing stale/messy records

Reducing lost opportunities

Improving sales visibility

Support Draft Agent

Preparing first responses

Human approval for risky cases

Better retention and trust

Reporting Agent

Summarizing weekly metrics

Detecting stalled issues

Better decisions

Content Ops Agent

Coordinating research/writing/review

Approval before publishing

More consistent output

Payment Onboarding Agent

Preparing delivery steps

Verifying events before access

Cleaner fulfillment

Why business framing matters

Production-ready clients do not buy agent theory. They buy saved time, reduced leakage, faster response, better visibility, fewer missed leads, cleaner delivery, and lower operational risk.

YK SYSTEMS Delivery Model

Layer

Deliverable

Client value

Discovery

Workflow map + bottleneck diagnosis

Shows what to fix first.

System design

Agent architecture + tool map

Turns chaos into a build plan.

Build

Automation + CRM + agent instructions

Creates working system.

Governance

Approvals + logs + permissions

Protects business trust.

Testing

Fake-data and edge-case tests

Proves system behavior.

Handoff

SOP + training + maintenance plan

Makes system usable after delivery.

Delivery thesis

YK SYSTEMS should sell the delivery layer between business workflows, tools, agents, automation, approvals, and handoff. The deliverable is

not a prompt. The deliverable is a governed operating process.

Implementation Path: Workflow Audit

Offer explanation

Find repeated work, lost-lead points, manual admin, missing records, and unclear ownership. Deliver a bottleneck map and fastest first

automation recommendation.

Deliverables

Current workflow review

Architecture map

Tool and data model

Build plan

Approval gates

Error and audit plan

Test checklist

Handoff notes

Implementation Path: Client Command Center Setup

Offer explanation

Build intake, CRM, owner alerts, confirmation emails, follow-up, AI summary, logs, and handoff. This is the strongest entry offer for small

businesses.

Deliverables

Current workflow review

Architecture map

Tool and data model

Build plan

Approval gates

Error and audit plan

Test checklist

Handoff notes

Implementation Path: AI Agent System Design

Offer explanation

Design agent roles, tools, permissions, memory, approval gates, risk levels, and test plan. This is for businesses ready to move past basic

automation.

Deliverables

Current workflow review

Architecture map

Tool and data model

Build plan

Approval gates

Error and audit plan

Test checklist

Handoff notes

Implementation Path: Multi-Agent Operating System

Offer explanation

Coordinate planner, researcher, operator, reviewer, reporting, CRM, and support agents with human approvals and audit logs. This is the

high-ceiling offer.

Deliverables

Current workflow review

Architecture map

Tool and data model

Build plan

Approval gates

Error and audit plan

Test checklist

Handoff notes

Implementation Path: Governance and Approval Layer

Offer explanation

Add policy, permissions, approvals, audit logs, incident response, and kill-switch logic to existing automations or agents.

Deliverables

Current workflow review

Architecture map

Tool and data model

Build plan

Approval gates

Error and audit plan

Test checklist

Handoff notes

Worksheet: Agent Role Canvas

Purpose

Agent name, business problem, user, goal, role, tone, inputs, outputs, tools, actions, approval rules, never actions, owner.

Fields to fill

Business owner

System name

Current state

Desired outcome

Required inputs

Tools/data sources

Approval rules

Error path

Test result

Next improvement

How to use it

Complete this worksheet before building. A filled worksheet is proof that the system has been designed, not improvised.

Worksheet: Agent Tool Permission Sheet

Purpose

Tool name, access type, read/write level, data touched, action risk, approval needed, credential owner, log requirement.

Fields to fill

- Business owner
- System name
- Current state
- Desired outcome
- Required inputs
- Tools/data sources
- Approval rules
- Error path
- Test result
- Next improvement

How to use it

Complete this worksheet before building. A filled worksheet is proof that the system has been designed, not improvised.

Worksheet: Agent Memory Plan

Purpose

Memory type, source, freshness, retention, owner, review cadence, deletion rule, contamination risk.

Fields to fill

- Business owner
- System name
- Current state
- Desired outcome
- Required inputs
- Tools/data sources
- Approval rules
- Error path
- Test result
- Next improvement

How to use it

Complete this worksheet before building. A filled worksheet is proof that the system has been designed, not improvised.

Worksheet: Approval Gate Map

Purpose

Action, risk level, required approver, approval channel, timeout, fallback action, log fields.

Fields to fill

- Business owner
- System name
- Current state
- Desired outcome
- Required inputs
- Tools/data sources
- Approval rules
- Error path
- Test result
- Next improvement

How to use it

Complete this worksheet before building. A filled worksheet is proof that the system has been designed, not improvised.

Worksheet: Agent Eval Scorecard

Purpose

Task, expected answer/action, pass/fail criteria, safety criteria, tool-use criteria, latency/cost target.

Fields to fill

- Business owner
- System name
- Current state
- Desired outcome
- Required inputs
- Tools/data sources
- Approval rules
- Error path
- Test result
- Next improvement

How to use it

Complete this worksheet before building. A filled worksheet is proof that the system has been designed, not improvised.

Worksheet: MCP Server Planning Sheet

Purpose

Server purpose, tools exposed, resources, authentication, transport, permissions, manifest review, threat model.

Fields to fill

- Business owner
- System name
- Current state
- Desired outcome
- Required inputs
- Tools/data sources
- Approval rules
- Error path
- Test result
- Next improvement

How to use it

Complete this worksheet before building. A filled worksheet is proof that the system has been designed, not improvised.

Worksheet: RAG Knowledge Base Planner

Purpose

Sources, chunking, metadata, update schedule, retrieval filters, citations, owner, stale-content policy.

Fields to fill

- Business owner
- System name
- Current state
- Desired outcome
- Required inputs
- Tools/data sources
- Approval rules
- Error path
- Test result
- Next improvement

How to use it

Complete this worksheet before building. A filled worksheet is proof that the system has been designed, not improvised.

Worksheet: Business OS Architecture Canvas

Purpose

Inputs, records, agents, automations, approvals, dashboards, logs, reports, owner routines.

Fields to fill

- Business owner
- System name
- Current state
- Desired outcome
- Required inputs
- Tools/data sources
- Approval rules
- Error path
- Test result
- Next improvement

How to use it

Complete this worksheet before building. A filled worksheet is proof that the system has been designed, not improvised.

Worksheet: Client Handoff Checklist

Purpose

System map, credentials without secrets, docs, SOP, test evidence, maintenance schedule, failure procedure.

Fields to fill

- Business owner
- System name
- Current state
- Desired outcome
- Required inputs
- Tools/data sources
- Approval rules
- Error path
- Test result
- Next improvement

How to use it

Complete this worksheet before building. A filled worksheet is proof that the system has been designed, not improvised.

Worksheet: Monthly Agent Audit

Purpose

Tool access, logs, failed actions, prompt changes, model changes, cost, security events, feedback, improvements.

Fields to fill

- Business owner
- System name
- Current state
- Desired outcome
- Required inputs
- Tools/data sources
- Approval rules
- Error path
- Test result
- Next improvement

How to use it

Complete this worksheet before building. A filled worksheet is proof that the system has been designed, not improvised.

Continue with the Deployment Toolkit

Use the companion Deployment Toolkit for editable templates, operating worksheets, mockups, tool starter kits, client delivery assets, and industry mini-packs.

Need help choosing the first workflow? Request a free automation / agent review at <https://hub.yksystems.ca>.